

Blind Catalytic Quantum Error Correction: Target-State Estimation and Fidelity Recovery Without A Priori Knowledge

Hikaru Wakaura and Taiki Tanimae

QIRI (Quantum Integrated Research Institute Inc.), 1-16-3, Akasaka, Minato-ku, Tokyo 107-0061, Japan
April 23, 2026

Near-term quantum computers must protect fragile coherence against decoherence to deliver useful results. Catalytic quantum error correction (CQEC) addresses this challenge by amplifying residual coherence with a reusable catalyst, achieving threshold-free recovery whenever the target coherent modes survive in the noisy state. However, the original protocol requires complete knowledge of the ideal target—an assumption that fails for variational and iterative algorithms whose output is unknown to the correction module. Here we show that this requirement can be removed by estimating the target from the noisy output alone, in a two-stage protocol we call *blind CQEC*. At the density-matrix level studied here, the recovery map reduces to a mode-inclusion-restricted projection of the estimate, so recovery fidelity is governed by estimation fidelity up to a mode-survival penalty; we make this structural property explicit and derive the corresponding bound $F_{\text{rec}} \geq 1 - 2\|\hat{\rho}_{\text{est}} - \rho_{\text{target}}\|_1 - \Delta_{\text{mode}}$, reducing the design of blind CQEC to a classical estimation problem bounded by a hard recovery ceiling. We benchmark five estimation strategies across three noise channels, four quantum algorithms ($d = 4$ –64), Haar-random states, and mixed targets, assuming oracle access to the noisy density matrix (the measurement cost of realizing this access is analyzed but not simulated). With exactly calibrated parameters, channel inversion coincides with the oracle on every benchmark, and both are capped by the mode-survival ceiling (down to $F_{\text{rec}} = 0.74$ at $d = 64$, and zero recoverability for the Regev state under strong dephasing); coherence maximization requires no noise model, matches the ceiling to within 1.5% at $d \leq 16$, and the operational choice between the two is set by calibration accuracy, with the dephasing rate exponentially dominant. A noisy-VQE demonstration for H_2 yields $3.4\times$ energy-error reduction, and a `qiskit-aer` exercise with oracle-calibrated noise parameters illustrates the pipeline on circuit-derived states. These results chart the estimator design space for blind, threshold-free recovery in algorithms where pre-encoding is unavailable, and identify the open problems — an operational measurement model and a circuit-level catalytic map — that separate this proof of concept from deployment.

1 Introduction

Quantum computers cannot deliver useful results unless the fragile coherence of their states is protected against decoherence [1]. The dominant remedy, stabilizer-based quantum error correction (QEC), works by encoding logical qubits into many physical qubits before noise acts and fails when the noise rate exceeds a code-specific threshold [2, 3, 4]. A complementary approach, catalytic quantum error correction (CQEC) [5], exploits the arbitrary amplification of quantum coherence under catalytic covariant transformations [6] to recover already-corrupted states without pre-encoding: recovery succeeds whenever the coherent modes of the target are present in

Hikaru Wakaura: h.wakaura@deeptell.jp

Taiki Tanimae: t.tanimae@deeptell.jp

the noisy state, $\mathcal{C}(\rho_{\text{target}}) \subseteq \mathcal{C}(\rho_{\text{noisy}})$, regardless of how small the residual coherence is, and the catalyst is reusable across unlimited correction cycles [5, 7]. This combination of threshold-free recovery and post-hoc applicability makes CQEC particularly attractive for near-term variational algorithms [8, 9] and iterative protocols where pre-encoding is impractical and the ideal output state evolves across rounds, leaving the error-correction module unable to reference a fixed target.

However, in its original formulation CQEC requires *complete knowledge of the target state* ρ_{target} , which has been identified as the principal conceptual gap between CQEC and conventional QEC [5].

In this study, we substantially relax this requirement by introducing *blind CQEC*,¹ where the target state ρ_{target} is estimated from the noisy state ρ_{noisy} (or multiple noisy copies) before catalytic correction. We found that the recovery fidelity is bounded as $F_{\text{rec}} \geq 1 - 2\|\hat{\rho}_{\text{est}} - \rho_{\text{target}}\|_1 - \Delta_{\text{mode}}$ (Sec. 2.4), where Δ_{mode} is the weight of target coherences whose modes the noise has annihilated, and numerically that the 84-condition benchmark realizes both regimes of this bound: recovery tracks estimation exactly where all modes survive ($r \geq 0.998$ at $d \leq 8$), and is capped by the mode-survival ceiling where they do not ($r = 0.77$ at $d = 64$, with a perfect estimate recovering only $F_{\text{rec}} = 0.41$ in the extreme case). The result shows that the performance of blind CQEC decomposes into a classical estimation problem plus a hard structural ceiling set by mode survival, reducing the design space of the protocol to the choice of estimator. We frame the problem as a two-stage protocol:

1. *Estimation*: Construct an estimate $\hat{\rho}_{\text{est}}$ from ρ_{noisy} and any available side information (noise model, copy count).
2. *Correction*: Apply standard CQEC with $\hat{\rho}_{\text{est}}$ as the proxy target.

The central question is: *how close can $F_{\text{rec}} = F(\rho_{\text{rec}}, \rho_{\text{target}})$ approach the oracle fidelity when the target is unknown?*

The nontrivial content of this question is twofold. First, it is not *a priori* clear that CQEC is robust to estimation error: the catalytic amplification protocol involves a nonlinear sequence of coherence transfers whose sensitivity to the target specification has not been analyzed. Second, while one could in principle apply any quantum state estimation method [10, 11, 12] before CQEC, the interplay between the estimation strategy and the subsequent catalytic recovery creates structure-specific constraints (e.g., the mode inclusion condition must hold for $\hat{\rho}_{\text{est}}$, not just ρ_{target}) that generic estimators do not address.

We developed five estimation strategies with varying information requirements (Table 1) and benchmarked them across four quantum algorithms from recent literature [13, 14, 15, 16], three noise channels, and Hilbert space dimensions $d = 4$ –64, characterizing the copy scaling of blind CQEC in the resource-constrained regime.

Our main results established three regimes:

- *Low-dimensional ($d \leq 16$), unknown noise model, near-pure target*. Coherence maximization, which needs no explicit noise model (though it implicitly favors phase-preserving channels), achieves $F_{\text{rec}} = 0.94$ –0.98, within 0.8–1.5% of the oracle ceiling; it must, however, be gated by a purity check, being harmful for target purity $v \lesssim 0.5$.
- *High-dimensional or mixed target, characterized noise*. Channel inversion coincides with the oracle under exact calibration ($F_{\text{rec}} = 0.74$ at $d = 64$, where the ceiling itself is set by mode survival), and retains its advantage only while the calibration error—dominated exponentially by the dephasing rate—stays within a narrow, asymmetric window.
- *Copy-constrained regime*. Within our fluctuation-averaging model, $n \approx 50$ –200 state preparations suffice for $F_{\text{rec}} \geq 0.95$ at $d \leq 16$; at $d = 64$ the threshold is unreachable at any copy count because the mode-survival ceiling lies below it.

¹We use “blind” in the operational sense that the target state is unknown to the correction protocol, distinct from blind quantum computing, which concerns hiding computation from an untrusted server.

2 Theoretical Background

2.1 Catalytic covariant transformations

We briefly recall the framework of Refs. [6, 5]. A covariant operation satisfies $\Lambda \circ \mathcal{U}_t = \mathcal{U}_t \circ \Lambda$ for all t , where $\mathcal{U}_t(\rho) = e^{-iHt}\rho e^{iHt}$. The modes of asymmetry $\mathcal{D}(\rho) = \{\Delta_{ij} = E_i - E_j \mid \rho_{ij} \neq 0\}$ characterize the coherence structure. The resonant coherent modes $\mathcal{C}(\rho) = \langle \mathcal{D}(\rho) \rangle_{\mathbb{Z}}$ form the integer lattice generated by $\mathcal{D}(\rho)$.

Theorem 1 (Shiraishi–Takagi [6]).—*A catalytic covariant transformation $\rho \otimes c \mapsto \tau$ with $\text{Tr}_C[\tau] = \rho'$, $\text{Tr}_S[\tau] = c$ exists if and only if $\mathcal{C}(\rho') \subseteq \mathcal{C}(\rho)$ and ρ is full rank.*

The CQEC protocol [5] applies Theorem 1 as follows: given a noisy state $\rho_{\text{noisy}} = \mathcal{E}(\rho_{\text{target}})$ and a catalyst c , find a covariant map Λ such that $\Lambda(\rho_{\text{noisy}} \otimes c) = \tau$ with $\text{Tr}_C[\tau] \approx \rho_{\text{target}}$. This requires knowing ρ_{target} explicitly to construct Λ . Note that Theorem 1 requires the input state to be full rank. In our benchmarks, this condition is satisfied because depolarizing noise ($\rho \mapsto (1-p)\rho + pI/d$) always produces full-rank states, and the combined noise model includes a depolarizing component; benchmarks without a depolarizing component rely on the mode-restriction step described below rather than on the full-rank guarantee.

2.2 Implementation via density-matrix simulation

We state precisely what is and is not simulated in this paper. The physical ICEC protocol of Ref. [5] constructs a catalyst c with $\mathcal{C}(c) \supseteq \mathcal{C}(\hat{\rho}_{\text{est}})$ and applies a catalytic covariant transformation to $\rho_{\text{noisy}} \otimes c$; simulating that transformation explicitly requires operating on the joint Hilbert space of dimension $d \cdot d_c$ (cost $O(d^6)$ per recovery for $d_c = d$), which is not done here. Instead, our benchmarks implement the *density-matrix-level interface* of the protocol, Eq. (1): the recovered state is the target estimate with (i) coherences restricted to the modes surviving in ρ_{noisy} (the mode-inclusion constraint of Theorem 1) and (ii) a PSD projection. All operations are $d \times d$ matrix manipulations of cost $O(d^3)$. This interface represents the *idealized* output of a perfectly executed catalytic transformation; it does not model the catalyst dynamics, gate compilation, decoherence during recovery, or any deviation of the physical protocol from its ideal action. Consequently, all fidelities reported in this paper are upper bounds on the performance of a physical implementation, and statements about “recovery” refer to this interface unless stated otherwise (Sec. 6.4).

2.3 The target-knowledge bottleneck

In standard CQEC, the target state enters in two places: (a) the mode inclusion check $\mathcal{C}(\rho_{\text{target}}) \subseteq \mathcal{C}(\rho_{\text{noisy}})$; (b) the construction of the catalytic map Λ , which steers amplification toward the off-diagonal structure of ρ_{target} . Without ρ_{target} , neither step can be performed exactly.

In blind CQEC, we replace ρ_{target} by an estimate $\hat{\rho}_{\text{est}}$ and accept a fidelity penalty. We make explicit at the outset a structural fact that shapes the interpretation of every numerical result in this paper: *at the density-matrix level used in our simulations, the recovery step reduces to a mode-inclusion-restricted projection of the estimate.* Concretely, the recovered state is

$$\rho_{\text{rec}} = \Pi(\mathcal{M}_{\rho_{\text{noisy}}}(\hat{\rho}_{\text{est}})), \quad (1)$$

where $\mathcal{M}_{\rho_{\text{noisy}}}$ zeroes those off-diagonal elements of $\hat{\rho}_{\text{est}}$ whose corresponding coherences are absent from ρ_{noisy} (the mode inclusion constraint of Theorem 1: coherent modes that have been completely destroyed cannot be amplified), and

$$\Pi(X) = \frac{X_+}{\text{Tr} X_+}, \quad X_+ \equiv \sum_i \max(0, \lambda_i) |v_i\rangle\langle v_i|, \quad (2)$$

is the projection onto the PSD cone followed by trace renormalization. The catalytic covariant transformation itself—the physical mechanism that would realize Eq. (1) on a device, including the explicit catalyst and the joint-system unitary—is *not* simulated here; its construction is the subject of Ref. [5], and the present paper takes Eq. (1) as the interface. An immediate consequence, developed below, is that recovery fidelity tracks estimation fidelity *by construction*: the near-perfect correlation reported in Sec. 5.3 is a verification of this structural property, not an independent empirical discovery about catalytic dynamics.

2.4 Recovery-fidelity bound

Because every estimator in this paper terminates with a PSD projection, the input $\hat{\rho}_{\text{est}}$ to the recovery step is itself a density matrix. In that case a rigorous bound follows directly from the Fuchs–van de Graaf inequality [17] ($1 - F(\rho, \sigma) \leq \|\rho - \sigma\|_1$ in the squared-fidelity convention of Eq. (10)) and the triangle inequality:

$$F_{\text{rec}} \geq 1 - \|\hat{\rho}_{\text{est}} - \rho_{\text{target}}\|_1 - \Delta_{\text{mode}}, \quad \Delta_{\text{mode}} \equiv \|\mathcal{M}_{\rho_{\text{noisy}}}(\hat{\rho}_{\text{est}}) - \hat{\rho}_{\text{est}}\|_1, \quad (3)$$

where Δ_{mode} is the mode-restriction penalty: the trace-norm weight of the estimated coherences that cannot be recovered because the corresponding modes are absent from ρ_{noisy} . Under the noise models of Sec. 4.2 at moderate strength, no coherence is driven exactly to zero and Δ_{mode} vanishes to numerical precision; it becomes appreciable only under deep dephasing at large mode gaps.

For general (not necessarily PSD) Hermitian inputs, one needs in addition the stability of the projection Π in trace norm. The clipping map $A \mapsto A_+$ is the metric projection onto the PSD cone with respect to the *Frobenius* norm and is therefore nonexpansive in that norm, $\|A_+ - B_+\|_F \leq \|A - B\|_F$, by the standard nonexpansiveness of projections onto closed convex sets. In trace norm the sharp constant is more delicate: the matrix absolute value (and hence the positive part) is not operator Lipschitz in general, and the best rigorous constants carry a logarithmic dimension dependence [18]. Numerically, over 10^5 random Hermitian pairs at $d \leq 16$ we observe $\|A_+ - B_+\|_1 \leq \|A - B\|_1$ with maximal ratio 1.000, and we conjecture that the constant-1 bound holds for trace-one Hermitian inputs; combining it with the renormalization step as in the display below then gives the working bound

$$F_{\text{rec}} \geq 1 - 2\|\hat{\rho}_{\text{est}} - \rho_{\text{target}}\|_1 - \Delta_{\text{mode}}, \quad (4)$$

via $\Pi(A) - \Pi(B) = \frac{1}{t_A}(A_+ - B_+) + (\frac{1}{t_A} - \frac{1}{t_B})B_+$ with $t_A = \text{Tr}A_+ \geq 1$ and $|t_A - t_B| \leq \|A - B\|_1$. We flag clearly which parts of this chain are proven and which are conjectured: Eq. (3) for density-matrix inputs is rigorous; Eq. (4) for general Hermitian inputs rests on the numerically supported constant-1 clipping bound (a rigorous version holds with an additional $O(\log d)$ factor [18]). We also caution that the empirical regression slope $a \approx 0.98$ of Sec. 5.3 is a descriptive statistic of the benchmark ensemble, not a measurement of a Lipschitz constant.

For estimators that consume n noisy state preparations, the estimation error decomposes as $\|\hat{\rho}_{\text{est}} - \rho_{\text{target}}\|_1 \leq \|\hat{\rho}_{\text{est}} - \rho_{\text{target}}\|_1 + O(n^{-1/2})$ into *systematic bias* and *statistical fluctuation*, giving

$$F_{\text{rec}} \geq 1 - 2(\|\hat{\rho}_{\text{est}} - \rho_{\text{target}}\|_1 + O(n^{-1/2})) - \Delta_{\text{mode}}. \quad (5)$$

Eqs. (3)–(5) motivate our approach: within this framework, minimizing the estimation error is the *only* available lever, which is why the remainder of the paper is devoted to estimator design. For circuit-level ICEC implementations [5], the catalytic amplification step introduces additional channel non-idealities that may further weaken L ; we leave the analytical extension to that setting for future work.

2.5 Dimension dependence of the estimation error

We give a scaling argument for why the two principal estimators degrade differently with dimension; the location of any crossover between them is determined empirically in Sec. 5.6 and depends on the accuracy of the assumed noise parameters.

For a Haar-random pure state $|\psi\rangle$ in dimension d , the expected populations scale as $\mathbb{E}[p_i] = 1/d$ with variance $\text{Var}[p_i] = (d-1)/[d^2(d+1)]$ [1], and the expected off-diagonal magnitudes obey $\mathbb{E}[|\rho_{ij}|^2] = 1/[d(d+1)]$ for $i \neq j$. Coherence maximization replaces each $|\rho_{ij}|$ by the physicality bound $\sqrt{p_i p_j} \approx 1/d$, so the relative per-element bias $|\sqrt{p_i p_j} - |\rho_{ij}||/|\rho_{ij}|$ does not vanish as d grows: the population spectrum carries progressively less information about the specific off-diagonal structure of the target, and the estimator degrades monotonically with d . This is the mechanism behind the monotonic decay of the coherence-maximization curve in the Haar-random sweep.

Channel inversion behaves differently: when the noise channel is invertible and its parameters are known *exactly*, the inversion is exact at every dimension and the estimation error is zero up

to numerical precision. Its practical limitation is therefore not dimension per se but *parameter misspecification*: a relative error δ in the assumed noise parameters produces a residual estimation error that grows with both δ and d (because the dephasing inversion factor $e^{+\gamma|\Delta_{ij}|}$ amplifies parameter errors exponentially in the mode gap). Consequently, the operationally relevant comparison is not “coherence maximization vs. exact inversion” but “coherence maximization vs. inversion at realistic calibration accuracy,” and the dimension at which the two cross depends on δ ; this trade-off surface is mapped quantitatively in the sensitivity analysis of Sec. 5.7.

An earlier version of this manuscript reported a closed-form crossover dimension obtained from a fixed-point equation; that equation admits no positive solution for our noise parameters and has been retracted. We thank the numerical audit accompanying this revision for identifying the error.

3 Estimation Strategies

This section presents the five estimation strategies, ordered by increasing information requirements.

3.1 Naive passthrough

The simplest approach uses $\hat{\rho}_{\text{est}} = \rho_{\text{noisy}}$ directly. Since ρ_{noisy} is the corrupted version of ρ_{target} , this is equivalent to running CQEC with the “wrong” target. As we show in Sec. 5, this yields $F_{\text{rec}} = F_{\text{noisy}}$ (no improvement), because catalytic amplification faithfully steers the state toward $\hat{\rho}_{\text{est}} = \rho_{\text{noisy}}$, which is itself the noisy state.

3.2 Coherence maximization

We propose maximizing the off-diagonal coherence of ρ_{noisy} while preserving its diagonal (population) structure:

$$(\hat{\rho}_{\text{est}})_{ij} = \begin{cases} (\rho_{\text{noisy}})_{ii} & \text{if } i = j, \\ \sqrt{(\rho_{\text{noisy}})_{ii}(\rho_{\text{noisy}})_{jj}} e^{i\phi_{ij}} & \text{if } i \neq j, \end{cases} \quad (6)$$

where $\phi_{ij} = \arg[(\rho_{\text{noisy}})_{ij}]$ when $(\rho_{\text{noisy}})_{ij} \neq 0$, and $\phi_{ij} = 0$ otherwise (coherences that have been completely suppressed cannot be recovered by this strategy; see below). This corresponds to constructing the *most coherent state consistent with the observed populations and phases*.

The rationale is physical: decoherence reduces $|\rho_{ij}|$ below the physicality bound $\sqrt{p_i p_j}$ but preserves the phase (for dephasing and depolarizing channels). By restoring $|\rho_{ij}|$ to its maximum allowed value, we undo the amplitude damping of coherences without requiring knowledge of the original amplitudes.

The matrix $\hat{\rho}_{\text{est}}$ constructed by Eq. (6) is Hermitian and trace-one by construction. It is positive semidefinite when the phases $\{\phi_{ij}\}$ are consistent with a rank-one state (i.e., $\phi_{ij} = \theta_i - \theta_j$ for some real $\{\theta_i\}$), which holds exactly for pure-state targets under phase-preserving noise. For general phase structures, the resulting matrix may have small negative eigenvalues; in such cases, we project onto the PSD cone via eigenvalue clipping (setting negative eigenvalues to zero and renormalizing). In our benchmarks, this projection is needed only for the $d = 64$ case, where the correction is $< 10^{-3}$ in trace norm.

For pure-state targets, $\hat{\rho}_{\text{est}}$ from Eq. (6) exactly recovers ρ_{target} when the noise is purely dephasing *and* no off-diagonal element is driven exactly to zero (i.e., $(\rho_{\text{noisy}})_{ij} \neq 0$ whenever

Table 1: Target-state estimation strategies for blind CQEC. The estimation complexity column refers to the cost of constructing $\hat{\rho}_{\text{est}}$ from ρ_{noisy} ; all strategies additionally incur $O(d^3)$ cost for the CQEC recovery step (matrix diagonalization and catalytic map construction).

Strategy	Info required	Est. complexity
Naive	None	$O(1)$
Coherence max	None	$O(d^2)$
Iterative	None	$O(k \cdot d^3)$
Multi-copy avg.	n copies	$O(nd^2)$
Channel inversion	Noise model	$O(d^2)$

($\rho_{\text{target}})_{ij} \neq 0$), since the populations are unchanged and the phases are preserved. In practice, finite dephasing rates produce exponentially suppressed but nonzero coherences, so this condition is satisfied for any finite γ ; however, numerical underflow may violate it for very large γ or $|\Delta_{ij}|$.

An important subtlety is the mode inclusion condition: Theorem 1 requires $\mathcal{C}(\hat{\rho}_{\text{est}}) \subseteq \mathcal{C}(\rho_{\text{noisy}})$. Since coherence maximization sets $(\hat{\rho}_{\text{est}})_{ij} = \sqrt{p_i p_j} e^{i\phi_{ij}}$, the estimate $\hat{\rho}_{\text{est}}$ has nonzero coherences wherever $p_i, p_j > 0$. For full-rank noisy states (guaranteed by depolarizing or combined noise), $\mathcal{C}(\rho_{\text{noisy}})$ includes all modes, so mode inclusion is automatically satisfied. For pure dephasing with $(\rho_{\text{noisy}})_{ij} = 0$ for some (i, j) , our convention $\phi_{ij} = 0$ ensures $(\hat{\rho}_{\text{est}})_{ij} = \sqrt{p_i p_j} \neq 0$, which may violate mode inclusion. In such cases, the ICEC protocol handles the violation by restricting amplification to the shared modes $\mathcal{C}(\hat{\rho}_{\text{est}}) \cap \mathcal{C}(\rho_{\text{noisy}})$, resulting in partial recovery.

We note that Eq. (6) can be viewed as a special case of constrained quantum state estimation [10], where the constraint set is the intersection of the PSD cone with the observed diagonal. Unlike general tomographic methods [11, 12], which reconstruct the state from informationally complete measurements, our estimator takes the noisy density matrix itself as input and exploits the specific structure of decoherence to reconstruct the off-diagonal elements. We stress that a density matrix is not a single measurement outcome: acquiring ρ_{noisy} operationally requires tomographic resources of order $\Theta(d^2/\epsilon^2)$ shots [12], and all “single state preparation” accounting in this paper is conditional on this oracle density-matrix access (see Secs. 5.10 and 5.11). This makes it fundamentally different from (and complementary to) standard QST approaches.

3.3 Noise-channel inversion

When the noise model $\mathcal{E}$ is known but the target is not, we invert the channel analytically:

$$\text{Dephasing: } (\hat{\rho}_{\text{est}})_{ij} = (\rho_{\text{noisy}})_{ij} e^{\gamma|\Delta_{ij}|}, \quad (7)$$

$$\text{Depolarizing: } \hat{\rho}_{\text{est}} = \frac{\rho_{\text{noisy}} - \frac{p}{d}I}{1-p}, \quad (8)$$

$$\text{Amp. damping: } (\hat{\rho}_{\text{est}})_{11} = \frac{(\rho_{\text{noisy}})_{11}}{1-\gamma}, \quad (\hat{\rho}_{\text{est}})_{01} = \frac{(\rho_{\text{noisy}})_{01}}{\sqrt{1-\gamma}}, \quad (9)$$

followed by projection onto the positive semidefinite cone (eigenvalue clipping with renormalization, applied to all three inversion formulas whenever the result is not PSD). Eq. (9) is written for a qubit; for $d > 2$, we generalize amplitude damping as cascaded decay $|k\rangle \rightarrow |k-1\rangle$ with rate γ per level, i.e., $(\rho_{\text{noisy}})_{jk} = (1-\gamma)^{(j+k)/2}(\rho_{\text{target}})_{jk}$ for $j, k > 0$, with population transfer to the ground state. The inversion formula generalizes accordingly: $(\hat{\rho}_{\text{est}})_{jk} = (\rho_{\text{noisy}})_{jk}/(1-\gamma)^{(j+k)/2}$. Channel inversion is exact when the channel is invertible (dephasing with $\gamma < \infty$, depolarizing with $p < 1$), but requires precise knowledge of the noise parameters. For depolarizing noise with large p (specifically $p > 1 - 1/d$), the inverted state from Eq. (8) may have negative eigenvalues even for a valid input ρ_{noisy} ; the PSD projection step corrects this at the cost of introducing estimation bias. In our benchmarks ($p = 0.3$, $d \leq 64$), PSD projection is rarely triggered for channel inversion.

3.4 Iterative refinement

Starting from the coherence-maximized estimate, we iterate:

1. Set $\hat{\rho}_{\text{est}}^{(0)}$ from Eq. (6).
2. For $k = 1, \dots, K$: run CQEC with target $\hat{\rho}_{\text{est}}^{(k-1)}$ to obtain $\rho_{\text{rec}}^{(k)}$; update $\hat{\rho}_{\text{est}}^{(k)} = \alpha \rho_{\text{rec}}^{(k)} + (1 - \alpha) \hat{\rho}_{\text{est}}^{(k-1)}$ with damping parameter $\alpha = 0.5$.

The damping parameter $\alpha = 0.5$ was selected empirically to balance convergence speed against oscillation; values $\alpha \in [0.3, 0.7]$ yield similar results in our benchmarks. We do not have an analytical convergence guarantee for this iteration: since the CQEC recovery map depends nonlinearly on the target estimate, the iteration is not a standard contraction mapping. In all cases tested ($d \leq 64$, $K \leq 20$), the iteration was observed to converge monotonically in fidelity within $K = 5$ steps, with $|F_{\text{est}}^{(k)} - F_{\text{est}}^{(k-1)}| < 10^{-4}$ serving as the stopping criterion. However, convergence failure at higher dimensions or under non-standard noise models cannot be excluded.

3.5 Multi-copy averaging

Given n_s independent noisy copies, we average $\bar{\rho} = n_s^{-1} \sum_{k=1}^{n_s} \rho_{\text{noisy}}^{(k)}$ (reducing statistical variance by $1/n_s$) and then apply coherence maximization. We emphasize that this procedure is *not* quantum state tomography in the standard sense [10, 12]: no informationally complete measurements are performed on the copies, and no maximum-likelihood or Bayesian estimation is used. Rather, the copies are assumed to be available as density matrices (consistent with our density-matrix-level simulation framework), and the averaging reduces fluctuations in the noisy state representation. In practice, the noisy copies are generated by independent applications of $\mathcal{E}$ to the unknown ρ_{target} , so $\bar{\rho}$ is a better estimate of $\mathcal{E}(\rho_{\text{target}})$, not of ρ_{target} directly. The subsequent coherence maximization step partially compensates for the systematic decoherence bias. Operationally, the n_s copies would be obtained by re-running the quantum algorithm n_s times (not by cloning, which is forbidden by the no-cloning theorem); the density-matrix averaging in our simulation is equivalent to the statistical average over independently prepared copies.

4 Benchmark Setup

4.1 Quantum algorithms

We tested blind CQEC on four quantum algorithms that produce specific output states subject to decoherence.

1. **qDRIFT** (Chen *et al.* [13]): Random product formula for Hamiltonian simulation of a 3-qubit Heisenberg chain ($d = 8$).
2. **QKAN** (Ivashkov *et al.* [14]): Quantum Kolmogorov–Arnold Network layer with Chebyshev degree $d_{\text{cheb}} = 3$, approximating $\sin(\pi x)$ ($d = 4$).
3. **Control-free QPE** (Clinton *et al.* [15]): Vectorial phase retrieval for spectral estimation of a 3-site Fermi–Hubbard model ($d = 16$).
4. **Regev factoring** (Regev [16]): Quantum factoring of $N = 15$ via discrete Gaussian state preparation and modular exponentiation ($d = 64$).

4.2 Noise models

Three noise channels are applied to the ideal algorithmic output. All benchmarks use the energy eigenbasis with equally spaced levels $E_i = i$ (the natural basis for the covariant framework of Sec. 2.1), so the mode gap is $|\Delta_{ij}| = |E_i - E_j| = |i - j|$.

1. *Dephasing* ($\gamma = 2.0$): $\rho_{ij} \mapsto e^{-\gamma|\Delta_{ij}|} \rho_{ij}$ for $i \neq j$. This corresponds to energy-gap-dependent dephasing, as arises naturally from coupling to a bosonic bath with spectral density proportional to the transition frequency [1]. Note that this is *not* equivalent to independent single-qubit dephasing; the gap-dependent model is chosen because the CQEC framework operates on energy modes Δ_{ij} , and coherences between distant energy levels decohere faster.
2. *Depolarizing* ($p = 0.3$): $\rho \mapsto (1 - p)\rho + \frac{p}{d}I$.
3. *Combined*: dephasing $\gamma = 1.0$, depolarizing $p = 0.15$, amplitude damping $\gamma_{\text{AD}} = 0.1$, applied sequentially in this order (dephasing first, then depolarizing, then amplitude damping). This composite model is not intended to represent a specific hardware platform but rather to stress-test the estimation strategies under a worst-case scenario where multiple decoherence mechanisms act simultaneously, as occurs generically in superconducting and trapped-ion systems (where T_1 relaxation, T_2 dephasing, and gate errors coexist).




Figure 1: **Coherence maximization is exact under pure dephasing.** Recovery fidelity (left), trace distance (center), and coherence ratio (right) versus dephasing strength γ for the five estimation strategies, the oracle (known target, dashed black), and the no-correction baseline. Coherence maximization and iterative refinement track the oracle exactly across $\gamma \in [0.1, 5.0]$, while naive estimation never improves on the uncorrected state.

4.3 Figures of merit

All fidelities are computed against the *true* ideal state ρ_{target} (not the estimated target):

$$F(\rho, \sigma) = \left(\text{Tr} \sqrt{\sqrt{\rho} \sigma \sqrt{\rho}} \right)^2. \quad (10)$$

We also report the trace distance $D(\rho_{\text{rec}}, \rho_{\text{target}}) = \frac{1}{2} \|\rho_{\text{rec}} - \rho_{\text{target}}\|_1$ and the coherence recovery ratio $C_{\ell_1}(\rho_{\text{rec}})/C_{\ell_1}(\rho_{\text{target}})$, where $C_{\ell_1}(\rho) = \sum_{i \neq j} |\rho_{ij}|$ is the ℓ_1 -norm of coherence [19], a standard monotone in the resource theory of coherence [20].

For the copy sweep, we fit $F(n) = 1 - A n^{-\alpha}$ to extract the scaling exponent α and the prefactor A . This power-law ansatz is motivated by the known $O(n^{-1})$ convergence of the infidelity $1 - F$ in standard quantum state tomography [12], where the exponent $\alpha = 1$ corresponds to the standard quantum limit. For our multi-copy averaging strategy (which is not full tomography), the effective exponent α may differ from unity, and we treat it as a free parameter to capture both sub-standard ($\alpha < 1$) and super-standard ($\alpha > 1$) convergence rates.

5 Results

5.1 Qubit noise sweeps

Figure 1 shows the recovery fidelity as a function of noise strength for a maximally coherent qubit ($d = 2$).

Under dephasing and depolarizing noise, coherence maximization, channel inversion, and iterative refinement all achieve $F_{\text{rec}} = 1.000$, matching the oracle exactly. The key insight is that these channels preserve the phase structure $\phi_{ij} = \arg(\rho_{ij})$ while attenuating only the magnitudes—exactly the information that coherence maximization restores.

Amplitude damping (Fig. 2) breaks this equivalence: only channel inversion achieves $F_{\text{rec}} = 0.9999$ (matching Oracle), while coherence maximization saturates at $F_{\text{rec}} = 0.926$ at $\gamma_{\text{AD}} = 0.5$. The failure is due to population redistribution ($\rho_{11} \rightarrow (1 - \gamma)\rho_{11}$), which cannot be inferred from the off-diagonal structure.

5.2 Four quantum algorithms

Table 2 presents the main results across all 4 algorithms under combined noise, with 5 estimation strategies plus no-correction baseline and oracle (all computed with the unified noise module and mode-restricted recovery; $n_{\text{copies}} = 100$).

Three features of Table 2 carry the physics. First, channel inversion equals the oracle to every reported digit on every algorithm—the consistency signature of an exact forward/inverse pair (Sec. 5.6). Second, coherence maximization, requiring no noise model, lands within 0.8–1.5% of




Figure 2: **Coherence maximization fails under amplitude damping; only channel inversion is exact.** Same axes as Fig. 1, with damping rate $\gamma_{AD} \in [0.05, 0.95]$. Channel inversion (blue) tracks the oracle; coherence maximization (green) saturates due to population redistribution that cannot be inferred from off-diagonal data alone.

Table 2: **Exactly calibrated channel inversion matches the oracle on every algorithm; coherence maximization follows within 1–15%; the Regev ceiling of 0.74 is the mode-inclusion constraint, not estimation error.** Recovery fidelity $F_{\text{rec}} = F(\rho_{\text{rec}}, \rho_{\text{target}})$ under combined noise ($\gamma = 1.0$, $p = 0.15$, $\gamma_{AD} = 0.1$). Bold values indicate the best blind strategy per algorithm.

Algorithm (d)	No corr.	Naive	Ch. inv.	Coh. max	Iter.	Multi-c.	Oracle
qDRIFT ($d = 8$)	0.195	0.195	0.955	0.942	0.849	0.523	0.955
QKAN ($d = 4$)	0.380	0.380	0.986	0.978	0.947	0.950	0.986
Control-free QPE ($d = 16$)	0.268	0.268	0.956	0.947	0.854	0.332	0.956
Regev factoring ($d = 64$)	0.369	0.369	0.738	0.586	0.586	0.048	0.738




Figure 3: **Channel inversion tracks the oracle in every cell; the visible dimension dependence is the recovery ceiling, not strategy ranking.** Heatmap of recovery fidelity across 4 algorithms $\times$ 3 noise models $\times$ 5 estimation strategies (plus no-correction baseline and oracle). Dark green indicates high fidelity; red indicates low fidelity.

the oracle at $d \leq 16$ and within 15% at $d = 64$. Third, the Regev ceiling of 0.738 (shared by the oracle) is a structural limit: under the combined noise only 44% of the Regev state’s coherence weight survives in recoverable modes, so no estimator—including the oracle—can exceed it.

To decompose the combined-noise results, Tables 3–5 present per-channel results for represen-

Table 3: Recovery fidelity under dephasing only ($\gamma = 2.0$). The Oracle column equals channel inversion at every entry: with exact parameters the inverted estimate is the true target, so the remaining deficit is the recovery-interface ceiling set by mode annihilation (coherences at large gaps fall below numerical support and cannot be recovered).

Strategy	$d=4$	$d=8$	$d=16$	$d=64$
No correction	0.414	0.280	0.117	0.036
Coherence max	0.942	0.828	0.622	0.148
Ch. inversion	0.942	0.828	0.622	0.523
Oracle	0.942	0.828	0.622	0.523

Table 4: Recovery fidelity under depolarizing only ($p = 0.3$). Depolarizing noise preserves the coherence-mode structure, so the recovery ceiling stays high at every dimension.

Strategy	$d=4$	$d=8$	$d=16$	$d=64$
No correction	0.775	0.738	0.719	0.705
Coherence max	0.943	0.975	0.968	0.952
Ch. inversion	0.998	0.992	0.984	0.973
Oracle	0.998	0.992	0.984	0.973

Table 5: Recovery fidelity under amplitude damping only ($\gamma_{\text{AD}} = 0.3$).

Strategy	$d=4$	$d=8$	$d=16$	$d=64$
No correction	0.815	0.400	0.173	0.040
Coherence max	0.966	0.869	0.618	0.440
Ch. inversion	0.998	0.962	0.709	0.528
Oracle	0.998	0.962	0.709	0.528

tative pure targets (one fixed Haar-random state per dimension), computed with the same unified noise module (Sec. 7). For the actual algorithm states the same run gives a sharper boundary case worth reporting: under dephasing-only noise ($\gamma = 2.0$) the Regev state’s coherence weight in surviving modes drops to *zero*, Theorem 1 forbids any recovery, and every strategy—oracle included—returns the no-correction fidelity (0.41). This is the cleanest illustration in our data set of the mode-inclusion constraint as a hard structural limit.

The per-channel decomposition supports three observations. (i) *The recovery-interface ceiling, not estimation, dominates at high dimension for gap-dependent channels*: under dephasing-only and amplitude-damping-only noise, even the oracle degrades with d (to 0.52–0.53 at $d = 64$) because coherences at large mode gaps are annihilated and, by Theorem 1, cannot be recovered by any estimator. (ii) *Under depolarizing noise the mode structure survives* and both the oracle ceiling and both estimators remain high at every dimension (≥ 0.95). (iii) *Coherence maximization matches exact channel inversion wherever the phase structure survives* (dephasing and depolarizing at $d \leq 16$) and falls behind where population information degrades (dephasing at $d = 64$: 0.148 vs. 0.523; amplitude damping at all d). With exactly calibrated parameters, channel inversion equals the oracle at every entry; its practical disadvantage appears only under parameter misspecification (Sec. 5.7).

5.3 Estimation–recovery correlation

Figure 4 plots estimation fidelity $F_{\text{est}} = F(\hat{\rho}_{\text{est}}, \rho_{\text{target}})$ against recovery fidelity F_{rec} across all 84 data points (4 algorithms $\times$ 3 noise models $\times$ 7 conditions: 5 strategies plus baseline and oracle).

We are explicit about what this figure does and does not show. Because the recovery interface is $\rho_{\text{rec}} = \Pi(\mathcal{M}_{\rho_{\text{noisy}}}(\hat{\rho}_{\text{est}}))$ [Eq. (1)], the relationship between the two fidelities is fixed by construction: $F_{\text{rec}} = F_{\text{est}}$ whenever the mode-restriction penalty Δ_{mode} vanishes, and F_{rec} falls below F_{est} by (at most) Δ_{mode} otherwise. The data conform to this structure quantitatively. At low dimension, where all modes survive, the points sit on the diagonal ($r = 0.999$ at $d = 4$, 0.998 at $d = 8$); as dimension grows and gap-dependent dephasing annihilates wide-gap coher-




Figure 4: **Recovery fidelity tracks estimation fidelity up to the mode-restriction penalty Δ_{mode} , exactly as Eq. (3) requires: on-diagonal where all modes survive, and below the diagonal where mode annihilation caps recovery.** Estimation fidelity vs. recovery fidelity for all 84 algorithm–noise–condition combinations (Pearson r : 0.999 at $d = 4$, 0.998 at $d = 8$, 0.930 at $d = 16$, 0.770 at $d = 64$; pooled 0.916). The largest deviation ($F_{\text{est}} = 1.00$ vs. $F_{\text{rec}} = 0.41$; Regev under dephasing-only) is a perfect estimate of a largely unrecoverable state.

ences, the correlation degrades exactly as the bound predicts ($r = 0.930$ at $d = 16$, 0.770 at $d = 64$), with the extreme case—a *perfect* estimate ($F_{\text{est}} = 1.00$) recovering only $F_{\text{rec}} = 0.41$ for the Regev state under dephasing-only noise—realizing the maximal mode-restriction penalty in our data set. The figure is therefore a quantitative verification of Eq. (3) in both of its regimes, not an independent empirical discovery about catalytic dynamics; within this framework, performance differences between strategies originate in the estimation step, and the recovery ceiling is set by mode survival. Whether a physical (circuit-level) catalytic implementation preserves this structure is an open question that cannot be answered at the density-matrix level; the pooled regression ($F_{\text{rec}} = 0.88 F_{\text{est}} + 0.03$, $R^2 = 0.84$; the 84 points are not i.i.d. and the statistics are descriptive) summarizes this benchmark ensemble only.

5.4 Copy scaling

Figure 5 shows F_{rec} versus copy count n_{copies} for each algorithm under combined noise.

Fitting $F(n) = 1 - A n^{-\alpha}$ via nonlinear least squares yielded the following scaling exponents. We list each exponent with the exact noise condition it was fitted under, since the exponents are condition-specific and must not be compared across different noise strengths:

- Channel inversion, depolarizing $p = 0.5$: $\alpha = 1.12$.




Figure 5: **Effective strategies at $d \leq 16$ reach $F_{\text{rec}} \geq 0.95$ with $n \approx 50\text{--}200$ copies; at $d = 64$ no strategy—including the oracle—reaches this threshold, because the mode-restricted ceiling (≈ 0.75) lies below it.** Fidelity vs. copy count n_{copies} for the four quantum algorithms under combined noise. Each panel shows the five blind strategies plus the no-correction baseline.

- Channel inversion, amplitude damping $\gamma_{\text{AD}} = 0.3$: $\alpha = 2.16$ (same condition, coherence max: $\alpha = 0.22$).
- Coherence max, dephasing $\gamma = 1.0$: $\alpha = 0.75$.
- Coherence max, amplitude damping $\gamma_{\text{AD}} = 0.7$: $\alpha = 0.04$ (same condition, channel inversion: $\alpha = 1.19$).

Two methodological caveats limit the physical interpretation of these exponents. First, the multi-copy ensemble is generated by adding fixed-amplitude Gaussian matrix perturbations ($\sigma = 0.01$ per element) to the noisy state, a proxy for state-preparation fluctuations that is *not* derived from Born-rule measurement statistics; the absolute perturbation scale means the effective signal-to-noise ratio varies with dimension. Second, the recovery convergence model contributes an n -dependence of its own that biases fitted exponents toward $\alpha \approx 1$ at large n . The exponents should therefore be read as descriptive summaries of this simulation pipeline; establishing measurement-derived scaling laws requires shot-level simulation with explicitly allocated POVMs, which we identify as necessary future work (Sec. 6.4). Within a fixed condition, the qualitative ordering is robust: strategies that address the dominant error mechanism gain from additional copies, while strategies that cannot correct the dominant error (coherence max under strong amplitude damping) show near-zero exponents regardless of copy count.

Given the methodological caveats above, we restrict physical interpretation to the within-condition contrast, which survives both caveats because they apply equally to all strategies within a condition: under amplitude damping at $\gamma_{\text{AD}} = 0.3$, channel inversion (which corrects the dominant error) gains rapidly with copies while coherence maximization (which cannot) gains slowly ($\alpha = 2.16$ vs. 0.22); at $\gamma_{\text{AD}} = 0.7$ the same ordering holds (1.19 vs. 0.04). Exponents above unity should not be over-interpreted as beating tomographic limits—they reflect the interaction of the perturbation model with the inversion step, not a measurement-theoretic speedup. (An earlier version of this manuscript included a figure interpreting these exponents against reference scal-

Table 6: **Minimum state preparations for $F_{\text{rec}} \geq 0.95$ per algorithm (combined noise); “ > 200 ” = threshold not reached.** “Preparations” counts noisy-state inputs under the fluctuation-averaging model of Sec. 5.4, conditional on oracle density-matrix access (Sec. 5.10); for strategies that consume one state representation (channel inversion, coherence maximization), averaging is applied as preprocessing. Channel inversion matches the oracle’s requirement in every column; Regev ($d = 64$) never reaches the threshold because the mode-restricted ceiling (≈ 0.75) lies below it.

Strategy	qDRIFT	QKAN	CF-QPE	Regev
Channel inversion	100	50	100	> 200
Coherence max	200	50	200	> 200
Iterative	> 200	200	> 200	> 200
Multi-copy avg.	> 200	50	> 200	> 200
Oracle	100	50	100	> 200




Figure 6: **Coherence maximization tracks the oracle up to $\gamma \approx 2.5$ for a qutrit.** Blind CQEC for a qutrit ($d = 3$) under dephasing, illustrating the intermediate behavior between qubit ($d = 2$) and the higher-dimensional algorithms.

ing laws; it has been withdrawn because the $n = 1$ point of the underlying sweep bypasses the fluctuation model entirely, making the fitted curves non-comparable across n .)

5.5 Dimension dependence

The qutrit ($d = 3$) benchmarks confirm intermediate behavior between qubit ($d = 2$) and the higher-dimensional algorithms. Figure 6 shows that coherence maximization maintains $F_{\text{rec}} > 0.95$ up to dephasing $\gamma \approx 2.0$ for the qutrit, while this threshold drops to $\gamma \approx 1.5$ for $d = 16$ (control-free QPE) and $\gamma \approx 0.5$ for $d = 64$ (Regev factoring). The degradation is monotonic in d , consistent with the loosening of the physicality bound $\sqrt{p_i p_j} \rightarrow 1/d$ as $d \rightarrow \infty$.

5.6 Dimension sweep with Haar-random states

The algorithm-specific benchmarks of Secs. 5.2–5.5 used fixed target states determined by each quantum algorithm. To assess strategy performance independent of target-state structure, we performed a systematic dimension sweep using Haar-random pure states (generated via `benchmark_blind_dimension_sweep.py`). For each dimension $d \in \{2, 4, 8, 16, 32, 64\}$, 20 Haar-random pure states were drawn, subjected to combined noise ($\gamma = 1.0$, $p = 0.15$, $\gamma_{\text{AD}} = 0.1$), and corrected using each estimation strategy. (Dimensions $d \geq 128$ are excluded because the exact dephasing inversion factor $e^{\gamma(d-1)}$ exceeds double-precision range; results previously reported at those dimensions were numerically invalid and have been withdrawn.)

Figure 7 shows the mean recovery fidelity as a function of d , and its structure is best read in two layers. First, *channel inversion with exact parameters coincides with the oracle to every reported digit at every dimension* (0.998, 0.973, 0.854, 0.634, 0.541, 0.519 for $d = 2$ –64): this is the expected signature of a consistent forward/inverse pair, since exact inversion reproduces the true target and the remaining deficit is the recovery-interface ceiling (mode annihilation under gap-dependent dephasing, cf. Table 3), not estimation error. Second, *coherence maximization tracks the ceiling closely at low dimension and separates from it as d grows* (0.979 $\rightarrow$ 0.191), with the gap to channel inversion opening from 0.02 at $d = 2$ to 0.33 at $d = 64$. Consequently, under exact calibration there is *no crossover*: channel inversion is never worse. The operationally meaningful crossover emerges only when calibration error is taken into account (Sec. 5.7): at $d = 64$, a -10% parameter error drops channel inversion to $F_{\text{rec}} = 0.12$, below coherence maximization’s noise-model-free 0.19.

The fidelity distributions across the 20 Haar-random targets concentrate with increasing d (standard deviations shrink from ≈ 0.02 at $d = 4$ to $\lesssim 0.01$ at $d = 64$), confirming that the reported means are representative (Fig. 8).

A key finding is the comparison with standard linear inversion tomography. Linear inversion—averaging n noisy copies and projecting onto the positive semidefinite cone—recovers an estimate of $\mathcal{E}(\rho_{\text{target}})$ rather than ρ_{target} itself. When this estimate is used as the CQEC target, the recovery fidelity matches the no-correction baseline (Fig. 9), because the systematic bias of the noise channel is not removed by averaging. This validates the necessity of decoherence-aware estimation strategies




Figure 7: **With exactly calibrated parameters, channel inversion coincides with the oracle at every dimension; both are capped by the recovery-interface ceiling, and coherence maximization falls progressively below them.** Recovery fidelity versus Hilbert space dimension for Haar-random pure states under combined noise. Each point is the mean over 20 Haar-random targets; shaded bands show $\pm 1\sigma$. Channel inversion (= oracle) declines from $F_{\text{rec}} = 0.998$ ($d = 2$) to 0.519 ($d = 64$) due to mode annihilation, while coherence maximization declines faster, from 0.979 to 0.191 .




Figure 8: **Concentration of measure tightens the fidelity distribution at large d while the gap between strategies widens.** Distribution of recovery fidelities for Haar-random states at selected dimensions.

(coherence maximization or channel inversion) for blind CQEC: generic quantum state tomography, which faithfully reconstructs the *noisy* state, is insufficient because CQEC requires an estimate of the *pre-noise* target.




Figure 9: **Linear-inversion tomography fails as a CQEC target estimator: it cannot substitute for decoherence-aware estimation.** Comparison of blind CQEC strategies with standard linear inversion tomography (QST baseline) at $d = 8$. Linear inversion (red dashed) performs at the no-correction level because it recovers $\mathcal{E}(\rho_{\text{target}})$ rather than ρ_{target} itself.

5.7 Noise-parameter sensitivity

Channel inversion (Sec. 3.3) requires knowledge of the noise parameters, which in practice must themselves be estimated. To quantify the robustness of channel inversion to parameter misspecification, we performed a systematic sensitivity analysis (via `benchmark_blind_sensitivity.py`).

Figure 10 shows a heatmap of F_{rec} versus (d, δ) , where δ is the relative error applied uniformly to all three noise parameters. At $d = 4$ the strategy is forgiving ($F_{\text{rec}} = 0.72$ at $\delta = -30\%$, ≥ 0.95 for $\delta \geq 0$), but the tolerance window narrows quickly with dimension: at $d = 64$ the exact-parameter value $F_{\text{rec}} = 0.52$ falls to 0.20 at $\delta = -5\%$, 0.12 at $\delta = -10\%$, and 0.34 at $\delta = +10\%$. The degradation is strongly *asymmetric*: underestimating the noise is far more damaging than overestimating it, because under-inversion leaves residual attenuation on exactly the wide-gap coherences that the mode-restriction step then removes permanently, whereas over-inversion errors are partially absorbed by the PSD projection. The operational consequence is a direct comparison with the noise-model-free alternative: at $d = 64$, channel inversion beats coherence maximization ($F_{\text{rec}} = 0.19$) only when the calibration error lies within roughly $[-5\%, +10\%]$; outside this window the noise-model-free strategy is preferable despite its lower ceiling.

Per-parameter analysis (Fig. 11) shows that the error budget is dominated entirely by the dephasing rate γ , which enters the inversion as $e^{+\gamma|\Delta_{ij}|}$ so that a relative error is amplified exponentially in the mode gap: at $d = 64$, a -10% error in γ alone reduces F_{rec} from 0.52 to 0.16, whereas the depolarizing probability p can be misspecified by $\pm 30\%$ with almost no effect (0.49–0.52) and the damping rate γ_{AD} is intermediate (0.22 at -30%). These results sharpen the calibration requirement considerably relative to a naive reading: at $d = 64$, randomized-benchmarking-grade characterization [21] must pin the *dephasing rate specifically* to within a few percent (with a bias toward overestimation) for channel inversion to retain its advantage; at $d \leq 8$ a $\pm 10\%$ tolerance suffices.

5.8 Mixed-state targets

All preceding benchmarks used pure or near-pure target states. To characterize blind CQEC on mixed-state targets, we generated Werner-like states $\rho(v) = v|\psi\rangle\langle\psi| + (1-v)I/d$ at $d = 8$ with purity parameter $v \in [0.3, 1.0]$ and 10 Haar-random $|\psi\rangle$ per v .

Figure 12 reveals that coherence maximization degrades for mixed targets: the physicality bound $\sqrt{p_i p_j}$ overshoots the true coherences, biasing the estimate. We report the operationally




Figure 10: **Channel inversion is sharply and asymmetrically sensitive to parameter misspecification: underestimating the noise is far more damaging than overestimating it, and the tolerance window narrows rapidly with dimension.** Recovery fidelity F_{rec} as a function of dimension d and relative parameter error δ (all three noise parameters perturbed by the same relative fraction). At $d = 4$, $F_{\text{rec}} \geq 0.72$ across $\delta \in [-30\%, +30\%]$; at $d = 64$, a -5% error already costs 61% of the fidelity ($0.52 \rightarrow 0.20$) while $+10\%$ costs 35% ($\rightarrow 0.34$).




Figure 11: **The dephasing rate γ dominates the sensitivity budget; the depolarizing probability p is nearly harmless to misspecify.** Per-parameter sensitivity at $d = 8$ (solid) and $d = 64$ (dashed), varying one parameter at a time. At $d = 64$: -10% error in γ alone gives $F_{\text{rec}} = 0.16$ (vs. 0.52 exact), while $\pm 30\%$ in p leaves F_{rec} essentially unchanged (0.49–0.52); γ_{AD} is intermediate.

critical consequence explicitly: for $v \lesssim 0.5$, coherence maximization is worse than applying no correction at all (at $v = 0.3$: corrected $F_{\text{rec}} = 0.78$ vs. no-correction 0.92), because the maximally-coherent ansatz actively injects coherence that the mixed target does not possess. Any practical deployment of coherence maximization therefore requires a purity check before application; absent one, the strategy can degrade the state it is meant to protect. Channel inversion, by contrast, is purity-robust: it dominates every other strategy at all purities tested ($F_{\text{rec}} = 0.996, 0.988, 0.961$,




Figure 12: **Coherence maximization becomes actively harmful for strongly mixed targets, while channel inversion is purity-robust.** Recovery fidelity versus target-state purity v for Werner-like states at $d = 8$ under combined noise. For $v \lesssim 0.5$, coherence maximization falls below the no-correction baseline: the maximally-coherent ansatz injects coherence the mixed target does not possess.

0.850 at $v = 0.3, 0.5, 0.8, 1.0$; the decline toward $v = 1$ reflects the recovery-interface ceiling at $d = 8$, not a mixed-state effect). This identifies mixed-state targets as a regime where noise-model knowledge is essential, even at low d , and qualifies the low-dimensional recommendation of Sec. 6.1.

5.9 Hybrid strategy

The complementary strengths of coherence maximization (noise-model-free, effective at low d) and channel inversion (noise-model-aware, effective at high d) motivate a hybrid estimator:

$$\hat{\rho}_{\text{est}}^{(\text{hyb})}(w) = w \hat{\rho}_{\text{est}}^{(\text{inv})} + (1-w) \hat{\rho}_{\text{est}}^{(\text{CM})}, \quad (11)$$

where $w \in [0, 1]$ interpolates between the two strategies.

Figure 13 shows F_{rec} versus w at each dimension under the unified noise implementation. The result is unambiguous: $F(w)$ is monotonically increasing in w at every dimension tested, so $w^* = 1$ (pure channel inversion) throughout, with no interior optimum. This corrects an earlier version of this analysis, in which an apparent interior optimum at intermediate dimensions was an artifact of an inconsistency between the forward noise and the inversion conventions; under the unified implementation, exact channel inversion reproduces the oracle (Sec. 5.6) and no admixture of coherence maximization can improve on it. The practical motivation for the hybrid family therefore rests entirely on *parameter uncertainty*: when the calibration error is large enough that channel inversion underperforms coherence maximization (Sec. 5.7), an intermediate w interpolates between the miscalibrated inversion and the noise-model-free fallback. Two caveats apply: the w^* reported here is selected by maximizing fidelity against the *true* target (an oracle analysis characterizing the achievable envelope, not an operational tuning rule), and a blind selection rule for w from calibration data alone remains future work.




Figure 13: **Under exact calibration the hybrid family is monotone in w with $w^* = 1$ at every dimension: exact channel inversion is never improved by admixing coherence maximization.** (a) Recovery fidelity vs. mixing weight w at selected dimensions under combined noise ($w = 0$: coherence maximization; $w = 1$: channel inversion). (b) Optimal weight vs. dimension. w^* is selected by maximizing fidelity against the *true* target (an oracle analysis; see text).




Figure 14: **Idealized density-matrix-level comparison: exact-inversion methods (PEC, blind ChInv) reconstruct the state; extrapolation- and purification-based idealizations (ZNE, VD) do not; blind CoM is the only noise-model-free method that does.** (a) Recovery fidelity F_{rec} vs. dimension for no correction, ZNE, VD, PEC, blind coherence maximization (CoM), and blind channel inversion (ChInv), all as density-matrix-level idealizations (see scope statement in text; the comparison is not resource-matched). (b) Bar chart at $d = 16$.

5.10 Comparison with standard quantum error mitigation

A natural question is how blind CQEC relates to established quantum error mitigation (QEM) methods for near-term devices: zero-noise extrapolation (ZNE) [22, 23], probabilistic error cancellation (PEC) [22, 24], and virtual distillation (VD) [25, 26]. We preface the comparison with an important scope statement: *the comparison below is not resource-matched*. The standard QEM methods are designed to correct expectation values of observables from finite measurement data, with well-characterized sampling overheads; here we evaluate density-matrix-level idealizations of each method (ZNE as element-wise Richardson extrapolation of the density matrix, PEC as exact channel inversion without quasiprobability sampling, VD as $\rho^2/\text{Tr}\rho^2$ without a measurement protocol) on the same Haar-random states under combined noise (15 states per dimension). The comparison therefore probes only one axis—what state-level information each idealized method can reconstruct—and says nothing about shot budgets, calibration cost, or bias-variance trade-offs, which would require a shot-level study.

Table 7: Classical post-processing cost and input requirements at $d = 8$ for the density-matrix-level idealizations (runtime is mean per recovery call in our simulation). “State preparations” counts the noisy-state inputs consumed per recovery *given oracle density-matrix access*; it excludes the measurement cost of acquiring the density matrix, which for tomographic acquisition scales as $\Theta(d^2/\epsilon^2)$ shots (Sec. 5.11) and dominates all entries in this table in any operational deployment.

Method	Runtime (ms)	State preps.	Noise info
Blind CoM	0.12	1	None
Blind ChInv	0.07	1	Full model
PEC (idealized)	0.05	1	Full model
VD (idealized)	0.01	2	None
ZNE (idealized)	1.23	3	None

Figure 14 supports three observations within this idealized setting:

- *Idealized PEC and blind ChInv are mathematically equivalent and both achieve $F_{\text{rec}} = 1.000$ at every dimension tested*, since with exactly known noise parameters both amount to exact channel inversion. We report this explicitly: in this idealization PEC loses nothing to blind ChInv; the distinction between the two lies entirely in the operational layer not simulated here (PEC reconstructs observables via quasiprobability sampling with overhead exponential in circuit volume, whereas a physical CQEC implementation would act on the state via a catalytic map whose cost is analyzed in Ref. [5]).
- *The ZNE and VD idealizations scale poorly with dimension* ($F_{\text{rec}} \approx 0.51$ and 0.67 at $d = 4$, collapsing to 0.03 – 0.04 at $d = 64$), reflecting that extrapolation and purification do not reconstruct the pre-noise state when noise strongly mixes populations. This is a statement about the state-reconstruction axis only—not a criticism of these methods on their design objective (expectation-value correction).
- *Blind CoM is the only method in this comparison that improves the state without noise-model knowledge*, reaching $F_{\text{rec}} > 0.95$ for $d \leq 16$; at $d = 64$ it degrades well below idealized PEC/ChInv (0.30 vs. 1.000), so its advantage is confined to the low-dimensional regime.

Table 7 summarizes the classical costs of the idealized pipelines. Within this scope, blind CoM occupies the niche of noise-model-free state-level improvement at low dimension; a resource-matched operational comparison (equal circuits, observables, total shots, and calibration budgets for all methods) is required before any deployment-level claim can be made, and we identify it as future work.

5.11 Sample-complexity context

The relevant information-theoretic reference point is the sample complexity of quantum state tomography: reconstructing an unknown d -dimensional state to trace-norm accuracy ϵ requires $\Theta(d^2/\epsilon^2)$ copies in general [12]. This bound frames the resource accounting of Sec. 5.10: because our simulations grant the estimators oracle access to the exact noisy density matrix, the copy counts reported in Table 6 quantify only the *state-preparation fluctuation-averaging* requirement within our perturbation model, not the measurement cost of acquiring ρ_{noisy} in the first place. A full operational analysis—allocating explicit POVMs and shot budgets to each estimator and comparing the resulting end-to-end sample complexity against the tomographic baseline—is a necessary next step that the density-matrix-level framework of this paper cannot provide. We therefore refrain from claims of saturating or beating tomographic rates; Fig. 15 shows the empirical infidelity-copy curves of the two effective strategies at $d = 8$ against the $O(n^{-1})$ and $O(n^{-1/2})$ reference slopes, with the reference lines drawn as visual guides (their prefactors are illustrative, not derived).

5.12 End-to-end application: VQE for H_2

To demonstrate practical utility, we applied blind CQEC within a variational quantum eigensolver (VQE) for the H_2 ground state energy at bond length $0.735 \text{ \AA}$ in the STO-3G basis ($d = 4$, two




Figure 15: **Empirical infidelity-copy curves at $d = 8$, shown against $O(n^{-1})$ and $O(n^{-1/2})$ reference slopes (visual guides with illustrative prefactors).** Because the copy ensemble derives from a Gaussian perturbation model rather than measurement statistics, these curves characterize the simulation pipeline; they do not establish saturation of information-theoretic bounds.




Figure 16: **Channel-inversion blind CQEC reduces the noisy-VQE energy error for H_2 by $3.4\times$ without target-state knowledge.** (a) Energy vs. iteration for noiseless (green), noisy without correction (red), blind CQEC with coherence maximization (blue), and blind CQEC with channel inversion (purple); dashed line marks the exact ground state energy. (b) Final energy error $|E_{\text{opt}} - E_{\text{exact}}|$ for each scenario.

qubits). At each VQE iteration, the ansatz state $|\psi(\theta)\rangle$ is subjected to combined noise ($\gamma = 0.5$, $p = 0.1$, $\gamma_{\text{AD}} = 0.05$), and the energy is evaluated as $\langle H \rangle = \text{Tr}(H\rho)$ on the (possibly corrected) state.

Table 8: Estimator exercise on qiskit-aer circuit-derived states: estimate fidelity against the ideal pure state for five small circuits under per-gate depolarizing noise ($p_{\text{depol}} = 0.10$). The channel-inversion parameter p_{eff} is *oracle-calibrated* from the trace overlap with the ideal state (see text); the CI column is therefore an upper baseline, not a blind result.

Test (d)	F_{noisy}	F_{CM}	F_{CI}	p_{eff}
GHZ (2q, $d=4$)	0.880	0.950	0.965	0.16
GHZ (3q, $d=8$)	0.804	0.880	0.926	0.22
W-like (2q, $d=4$)	0.822	0.984	0.970	0.24
Random (2q, $d=4$)	0.746	0.885	0.972	0.34
Random (3q, $d=8$)	0.615	0.830	0.877	0.44

Figure 16 compares four scenarios: noiseless (ideal), noisy without correction, blind CQEC with coherence maximization, and blind CQEC with channel inversion. The noiseless VQE converges to the exact ground state ($E_0 = -1.851$ Ha). Without correction, noise induces an energy bias of $\Delta E = 0.34$ Ha. Channel-inversion blind CQEC reduces this to $\Delta E = 0.10$ Ha—a $3.4\times$ improvement—without any knowledge of the ansatz state at each iteration. Coherence maximization provides a marginal improvement ($\Delta E = 0.32$ Ha), reflecting the fact that VQE noise is dominated by amplitude damping (population relaxation), which coherence maximization cannot address (Sec. 3.2).

This demonstrates that blind CQEC can provide tangible benefits in a quantum chemistry workflow, even when the target state evolves across optimization iterations and is never explicitly known.

For balance we also report a negative result from the same benchmark suite: an analogous VQE run for LiH at $d = 16$ failed to converge in *all four* scenarios—including the noiseless one—with final energy errors of ≈ 6.2 Ha throughout, so blind CQEC neither helped nor harmed (the optimizer, not decoherence, was the bottleneck). The H_2 result should therefore be read as a proof of principle on a case where the VQE itself converges, not as evidence that blind CQEC rescues arbitrary variational workflows; characterizing when correction helps in larger active spaces requires a systematically tuned optimizer study.

5.13 Circuit-derived states: an estimator exercise with oracle calibration

All preceding results operate on states produced directly by our analytical noise channels. As a complementary exercise, we applied the estimation strategies to density matrices produced by the qiskit-aer simulator (v0.17, with qiskit v2.4) for five small state-preparation circuits: GHZ states on $n = 2$ and $n = 3$ qubits, a W-like state, and two Haar-random parameterized circuits (benchmark_blind_circuit_sanity.py), each executed under a per-gate depolarizing noise model with $p_{\text{depol}} = 0.10$ on single- and two-qubit gates.

Two scope limitations must be stated plainly, because they mean this exercise is *not* a blind circuit-level validation. First, the effective global depolarizing rate p_{eff} used by channel inversion is calibrated from the trace overlap with the *ideal* circuit output—information a blind protocol would not possess; in a deployment, p_{eff} would come from independent calibration (e.g., randomized benchmarking [21]), and the values reported here are oracle-calibrated upper baselines. Second, the reported fidelities are those of the *estimates* against the ideal state; consistent with the recovery interface of Eq. (1), no catalytic circuit is executed. What the exercise does show is that the estimators behave on circuit-derived noisy states (whose noise is a gate-local process, not our analytical channel) the same way they behave on channel-generated states.

Table 8 reports the resulting fidelities. Both estimators improve the estimate over the no-correction baseline in every case: coherence maximization (fully blind) gains 5–22%, and oracle-calibrated channel inversion gains 9–26%. The qualitative ordering of strategies is consistent with the density-matrix-level results at $d = 4, 8$ (Table 4). A genuine circuit-level validation of blind CQEC—with independently calibrated noise parameters and an executed catalytic recovery circuit—remains open and is listed in Sec. 6.4.




Figure 17: **Estimator behavior on circuit-derived noisy states matches the channel-generated benchmarks; channel inversion here uses oracle-calibrated noise parameters.** Five small state-preparation circuits under per-gate depolarizing noise ($p_{\text{depol}} = 0.10$) on qiskit-aer. Both estimators improve over the no-correction baseline in every case; the CI bars are oracle-calibrated upper baselines (see text).

6 Discussion

6.1 Regimes of applicability

Our results delineate three operational regimes for blind CQEC:

1. **Low-dim, unknown noise, near-pure target** ($d \leq 16$, purity $v \gtrsim 0.7$): Coherence maximization is sufficient. No noise-model information is needed; $F_{\text{rec}} > 0.95$ with ~ 10 copies. The purity qualifier is essential: below $v \approx 0.7$ coherence maximization is *worse than no correction* (Sec. 5.8), so a purity check must gate its use. This regime covers most near-term quantum algorithms with nearly pure outputs (qDRIFT, QKAN, QPE).
2. **High-dim or mixed target, characterized noise**: Channel inversion is required. Coherence maximization fails because population estimates become unreliable at high d and because the physicality bound overshoots for mixed targets. This regime requires noise characterization (e.g., via randomized benchmarking [21]) at the accuracy quantified in Sec. 5.7, but not target-state knowledge.
3. **Copy-constrained**: With only $n \leq 5$ copies, all blind strategies degrade significantly. The oracle itself achieves only $F_{\text{rec}} \approx 0.7$ at $n = 2$, indicating that the finite-copy penalty is intrinsic to the recovery interface, not to blindness.

A notable feature of the results is that *naive estimation never improves over no correction*. This occurs because CQEC is a directed recovery process that amplifies coherence modes selectively toward the target specification: when the target is set to the noisy state itself ($\hat{\rho}_{\text{est}} = \rho_{\text{noisy}}$), the noisy state is a fixed point of the recovery map, so no improvement occurs. Multi-copy averaging of identical noisy copies likewise yields ρ_{noisy} , so the naive strategy cannot be rescued by additional copies.

6.2 Why coherence maximization degrades at high dimension

The success of coherence maximization at low d and its degradation at high d can be understood quantitatively. For a d -dimensional maximally mixed noisy state, all populations approach $1/d$, and the physicality bound $\sqrt{p_i p_j} \approx 1/d$ carries vanishing information about the target's

off-diagonal structure. In contrast, at low d with near-pure targets, the bound is tight and phase-preserving noise retains the exact phase structure (Sec. 3.2), allowing near-perfect reconstruction. The widening gap between coherence maximization and the recovery ceiling as d grows thus reflects the transition from the information-rich to the information-poor regime of the population-based estimator; whether this translates into a preference for channel inversion in practice depends on the available calibration accuracy (Sec. 5.7).

6.3 Comparison with standard and purification QEC

It is important to position blind CQEC relative to standard stabilizer-based QEC [2, 3]. Standard QEC encodes logical qubits into physical qubits *before* noise acts, enabling syndrome-based correction without knowledge of the encoded state. This encoding-first paradigm is fundamentally different from CQEC, which operates *after* noise has corrupted an unencoded state. The two approaches are therefore not directly comparable: standard QEC requires foresight (pre-encoding), while blind CQEC is a post-hoc recovery method applicable when encoding was not performed or was insufficient. Blind CQEC does not replace standard QEC but rather addresses a complementary scenario—recovering states from noisy quantum computations where pre-encoding was absent or where the noise exceeded the code’s correction capacity.

Purification QEC (PQEC) [27] (arXiv:2603.11568) also recovers states without target knowledge, using recursive swap tests. The key structural differences are:

- PQEC uses recursive purification requiring multiple identical noisy copies, with the copy count growing with the desired precision; blind CQEC requires as few as 5–10 copies at $d \leq 16$.
- PQEC operates without any noise-model information; blind CQEC performs best when the noise model is known (channel inversion), though coherence maximization provides a noise-model-free alternative at low d .

The two approaches address complementary regimes: PQEC is suited for settings where no noise information is available and copies are abundant; blind CQEC is suited for the copy-constrained, noise-characterized regime.

6.4 Limitations

Several limitations warrant discussion; the first two are structural and bound the scope of every claim in this paper.

1. *The catalytic transformation is not simulated.* As stated in Secs. 2.2 and 5.3, all benchmarks implement the density-matrix-level recovery interface [Eq. (1)], which represents the idealized output of a perfectly executed catalytic covariant transformation. The explicit CPTP map acting on $\rho_{\text{noisy}} \otimes c$ —with verifiable catalyst preservation, covariance, and mode-inclusion checks—is constructed in Ref. [5] but is not integrated into these benchmarks; consequently, all reported fidelities are upper bounds, and properties such as the estimation–recovery transparency of Sec. 5.3 are properties of the interface, not verified properties of catalytic dynamics.
2. *No operational measurement model.* All estimators consume the exact noisy density matrix, which is not a physically accessible object at unit cost: acquiring it requires $\Theta(d^2/\epsilon^2)$ measurement shots. Every resource count in this paper is conditional on this oracle access (Secs. 5.10 and 5.11), and the multi-copy ensembles use a Gaussian perturbation proxy rather than Born-rule statistics (Sec. 5.4).
3. *Noise-parameter sensitivity of channel inversion.* The sensitivity analysis of Sec. 5.7 shows a narrow, asymmetric tolerance window at high dimension: at $d = 64$, the advantage over the noise-model-free strategy survives only for calibration errors within roughly $[-5\%, +10\%]$, with the dephasing rate as the dominant (exponentially amplified) parameter; at $d \leq 8$, $\pm 10\%$ suffices.

4. *Mixed-state targets.* Coherence maximization becomes actively harmful below purity $v \approx 0.5$ (Sec. 5.8) and must be gated by a purity check; channel inversion is purity-robust.
5. *Scalability and tensor product structure.* The Haar-random sweep covers $d \leq 64$ as a single d -level system ($d \geq 128$ exceeds double-precision range for the exact dephasing inversion). In realistic multi-qubit settings, noise acts locally, and estimation strategies could exploit this locality (e.g., per-qubit coherence maximization); extending the framework to tensor-product structure is a natural next step.
6. *Experimental validation.* All results are numerical; circuit-level implementation with independently calibrated noise and hardware validation remain open (the exercise of Sec. 5.13 uses oracle-calibrated parameters and is not a substitute).
7. *Dependence on companion preprints.* The unified CQEC and ICEC framework is given in our companion preprint [5] (arXiv:2603.25774), the catalyst-preparation methods in [7] (Research Square, doi:10.21203/rs.3.rs-9366084), and the PQEC comparison in [27] (arXiv:2603.11568). The bounds Eqs. (3)–(4) and the estimation strategies are self-contained.

7 Computational Details

All simulations were performed as density-matrix-level numerical calculations using custom Python code (NumPy 2.0.2, SciPy 1.13.1) on an Apple M4 Max workstation. No quantum hardware was used; apart from the `qiskit-aer` exercise of Sec. 5.13, all operations act directly on $d \times d$ density matrices at cost $O(d^3)$ per recovery (Sec. 2.2); the full catalytic joint-system simulation, not performed here, would cost $O(d^6)$.

Unified noise module and numerical audit. All noise channels (forward and inverse) are implemented once, in `blind_cqec/noise.py` and `blind_cqec/estimators.py`, and every benchmark script imports or reproduces these implementations; unit tests verify element-wise agreement between the scripts and the package (maximal deviation $\sim 10^{-16}$) and round-trip exactness of the inversion ($F > 0.999$ for $\mathcal{E}^{-1}(\mathcal{E}(\rho))$ at all $d \leq 64$). The mode-inclusion threshold is fixed at 10^{-10} (matching the mode checker of the core library): coherences of ρ_{noisy} below this magnitude are treated as annihilated and are zeroed in the recovery target [Eq. (1)]; this convention determines the $d = 64$ ceilings reported in Secs. 5.2 and 5.6. This revision followed an adversarial numerical audit that identified, in an earlier version, an inconsistency between the forward dephasing convention of one benchmark script and the gap-dependent inversion (which invalidated the previously reported dimension-crossover, sensitivity, and hybrid results), a numerically overflowing regime at $d \geq 128$, and a fixed-point crossover equation with no positive solution. All affected tables, figures, and claims have been regenerated from the unified module or withdrawn, as noted in the text.

All results are deterministic given the recorded seeds. The target states are fixed by the choice of quantum algorithm (not randomly sampled) in the algorithm benchmarks, and Haar-random ensembles (20 states per dimension) are used for the dimension sweep, sensitivity, mixed-state, and hybrid benchmarks, with means and standard deviations reported.

The complete source code is publicly available at https://github.com/deeptell-inc/blind_cqec_pkg, including the reference implementation as a Python package (`blind_cqec`, also published on PyPI) and all benchmark scripts:

- `benchmark_blind_cqec.py` – qubit/qutrit noise sweeps;
- `benchmark_blind_dimension_sweep.py` – Haar-random dimension sweep;
- `benchmark_blind_sensitivity.py` – noise-parameter sensitivity;
- `benchmark_blind_circuit_sanity.py` – qiskit-aer circuit-level sanity check;
- `benchmark_blind_vqe_demo.py` – end-to-end VQE demonstration for H_2 ;
- `benchmark_blind_algorithms.py`, `benchmark_blind_qem_compare.py`, and the figure-generation scripts.

8 Conclusion

We introduced blind catalytic quantum error correction: a two-stage protocol that removes the target-state knowledge requirement of CQEC by combining a classical pre-estimator with the density-matrix-level recovery interface of Ref. [5]. We stated at the outset the structural fact that governs every result: at this level the recovery map reduces to a mode-inclusion-restricted projection of the estimate [Eq. (1)], so recovery fidelity is bounded by estimation fidelity by construction ($F_{\text{rec}} \geq 1 - \|\hat{\rho}_{\text{est}} - \rho_{\text{target}}\|_1 - \Delta_{\text{mode}}$ for density-matrix estimates, rigorously; Eq. (4) in the general case). The 84-condition benchmark confirms this bound is respected everywhere, reducing the design of blind CQEC—within this framework—to a classical estimation problem.

The benchmark’s substantive content is the map of the estimator design space, which the unified noise implementation makes sharp. With exactly calibrated parameters, channel inversion coincides with the oracle at every dimension, and the residual deficit ($0.998 \rightarrow 0.519$ over $d = 2\text{--}64$) is the recovery-interface ceiling set by mode annihilation, not estimation error; the noise-model-free alternative, coherence maximization, tracks the ceiling at low dimension and separates from it as populations approach $1/d$ ($0.979 \rightarrow 0.191$). The operationally decisive variable is calibration accuracy: at $d = 64$, a -10% error in the noise parameters (dominated entirely by the dephasing rate) drops channel inversion below the noise-model-free strategy, and the sensitivity surface of Sec. 5.7 quantifies exactly when each estimator should be preferred. Two failure modes we report with equal prominence: coherence maximization is actively harmful for target purity $v \lesssim 0.5$, and no blind strategy reaches $F_{\text{rec}} = 0.95$ at $d = 64$ under combined noise.

These findings position blind CQEC as a candidate complement to the established error-handling stack—post-hoc and threshold-free where QEC is encoding-first and threshold-bounded, state-level where mitigation is observable-level—while the limitations analysis makes the gap to deployment explicit: the catalytic transformation itself is not simulated here, density-matrix access stands in for a measurement model whose tomographic cost dominates all resource counts, and the H_2 demonstration ($3.4\times$ energy-error reduction, Fig. 16) is a proof of principle accompanied by a null result on LiH. The path from this proof of concept to a deployable protocol runs through three open problems, in order: an explicit CPTP catalytic recovery map evaluated on $\rho_{\text{noisy}} \otimes c$ (the construction of Ref. [5], integrated end-to-end); an operational measurement model with explicit POVMs and shot budgets; and circuit-level validation with independently calibrated noise. The estimator design-space map presented here—exact-inversion ceiling, noise-model-free fallback, calibration-accuracy crossover, and purity gating—is the part of the problem that is now settled at the interface level and will carry over to any such implementation.

Acknowledgments

Numerical simulations were performed using the scientific Python ecosystem (NumPy, SciPy, Matplotlib) on Apple Silicon hardware. The reference implementation `blind_cqec` is released as an open-source Python package on PyPI under the MIT license.

Data and code availability

All simulation code, raw numerical data, and plotting scripts used to generate every figure and table in this paper are publicly available at https://github.com/deeptell-inc/blind_cqec_pkg. All results are deterministically reproducible with the fixed random seed `numpy.random.seed(42)`; the test suite shipped with the `blind_cqec` package locks in the headline numbers reported here.

References

- [1] Michael A. Nielsen and Isaac L. Chuang. “Quantum computation and quantum information”. Cambridge University Press. (2010). 10th anniversary edition.
- [2] Peter W. Shor. “Scheme for reducing decoherence in quantum computer memory”. *Phys. Rev. A* **52**, R2493 (1995).

- [3] Daniel Gottesman. “Stabilizer codes and quantum error correction”. Ph.D. thesis, Caltech (1997). [arXiv:quant-ph/9705052](#).
- [4] Emanuel Knill. “Quantum computing with realistically noisy devices”. *Nature* **434**, 39–44 (2005).
- [5] Hikaru Wakaura and Taiki Tanimae. “Catalytic quantum error correction: Theory, efficient catalyst preparation, and numerical benchmarks” (2026). [arXiv:2603.25774v4](#).
- [6] Naoto Shiraishi and Ryuji Takagi. “Arbitrary amplification of quantum coherence in asymptotic and catalytic transformation”. *Phys. Rev. Lett.* **132**, 180202 (2024).
- [7] Hikaru Wakaura and Taiki Tanimae. “Resource-efficient catalyst preparation for catalytic quantum state recovery via dynamical decoupling, twirling, and purification”. *Research Square preprint* (2026).
- [8] Alberto Peruzzo, Jarrod McClean, Peter Shadbolt, Man-Hong Yung, Xiao-Qi Zhou, Peter J. Love, Alán Aspuru-Guzik, and Jeremy L. O’Brien. “A variational eigenvalue solver on a photonic quantum processor”. *Nat. Commun.* **5**, 4213 (2014).
- [9] M. Cerezo, Andrew Arrasmith, Ryan Babbush, Simon C. Benjamin, Suguru Endo, Keisuke Fujii, Jarrod R. McClean, Kosuke Mitarai, Xiao Yuan, Lukasz Cincio, and Patrick J. Coles. “Variational quantum algorithms”. *Nat. Rev. Phys.* **3**, 625–644 (2021).
- [10] Zdeněk Hradil. “Quantum-state estimation”. *Phys. Rev. A* **55**, R1561 (1997).
- [11] David Gross, Yi-Kai Liu, Steven T. Flammia, Stephen Becker, and Jens Eisert. “Quantum state tomography via compressed sensing”. *Phys. Rev. Lett.* **105**, 150401 (2010).
- [12] Jeongwan Haah, Arlenn W. Harrow, Zhengfeng Ji, Xiaodi Wu, and Nengkun Yu. “Sample-optimal tomography of quantum states”. *IEEE Trans. Inf. Theory* **63**, 5628–5641 (2017).
- [13] Chi-Fang Chen, Hsin-Yuan Huang, Richard Kueng, and Joel A. Tropp. “Concentration for random product formulas”. *PRX Quantum* **2**, 040305 (2021).
- [14] Petr Ivashkov, Po-Wei Huang, Kelvin Koor, Lirandë Pira, and Patrick Rebentrost. “QKAN: Quantum Kolmogorov–Arnold networks” (2024). [arXiv:2410.04435](#).
- [15] Laura Clinton, Toby S. Cubitt, Raul Garcia-Patron, Ashley Montanaro, Stasja Stanisic, and Maarten Stroeck. “Quantum phase estimation without controlled unitaries”. *PRX Quantum* **7**, 010345 (2026).
- [16] Oded Regev. “An efficient quantum factoring algorithm” (2024). [arXiv:2308.06572v3](#).
- [17] Christopher A. Fuchs and Jeroen van de Graaf. “Cryptographic distinguishability measures for quantum-mechanical states”. *IEEE Trans. Inf. Theory* **45**, 1216–1227 (1999).
- [18] E. Brian Davies. “Lipschitz continuity of functions of operators in the Schatten classes”. *J. London Math. Soc.* **s2-37**, 148–157 (1988).
- [19] Tillmann Baumgratz, Marcus Cramer, and Martin B. Plenio. “Quantifying coherence”. *Phys. Rev. Lett.* **113**, 140401 (2014).
- [20] Alexander Streltsov, Gerardo Adesso, and Martin B. Plenio. “Colloquium: Quantum coherence as a resource”. *Rev. Mod. Phys.* **89**, 041003 (2017).
- [21] Joseph Emerson, Robert Alicki, and Karol Życzkowski. “Scalable noise estimation with random unitary operators”. *J. Opt. B* **7**, S347 (2005).
- [22] Kristan Temme, Sergey Bravyi, and Jay M. Gambetta. “Error mitigation for short-depth quantum circuits”. *Phys. Rev. Lett.* **119**, 180509 (2017).
- [23] Ying Li and Simon C. Benjamin. “Efficient variational quantum simulator incorporating active error minimization”. *Phys. Rev. X* **7**, 021050 (2017).
- [24] Suguru Endo, Simon C. Benjamin, and Ying Li. “Practical quantum error mitigation for near-future applications”. *Phys. Rev. X* **8**, 031027 (2018).
- [25] Bálint Koczor. “Exponential error suppression for near-term quantum devices”. *Phys. Rev. X* **11**, 031057 (2021).

- [26] William J. Huggins, Sam McArdle, Thomas E. O’Brien, Joonho Lee, Nicholas C. Rubin, Sergio Boixo, K. Birgitta Whaley, Ryan Babbush, and Jarrod R. McClean. “Virtual distillation for quantum error mitigation”. *Phys. Rev. X* **11**, 041036 (2021).
- [27] Rounak Raghoonanan and Tim Byrnes. “Purification quantum error correction” (2026). [arXiv:2603.11568v1](#).